

Exercise: Distributional Analysis Using Pandas

Summary

This exercise uses Pandas to do the ETR calculation from the earlier assignment.

Input Data

Files **households.csv** and **quantities.csv** are the CSV files from the previous distributional analysis. As you'll probably recall, there are 1000 households in the analysis and households.csv has their attributes. It has five columns, **id**, **type**, **inc**, **a** and **b**, and there is one row for each household. The second file, quantities.csv, has three columns: **id**, **qd1**, and **qd2** and again there is one row for each household.

Deliverables

A script called **etr.py** that carries out the calculations described below.

Instructions

1. Import pandas as `pd`.
2. Define a function called `print_groups()` that takes two arguments, a dataframe of ETRs called `hh`, and a list of variables to be used to group the data called `group_vars`. The body of the function should do the following:
 1. Group the data by setting variable `grouped` to the result of calling the `.groupby()` method on `hh` using `group_vars` as the argument.
 2. Set variable `med` to the result of calling the `.median()` method on `grouped` column "etr".
 3. Round the results to two digits by setting variable `med` to the result of calling the `.round()` method of `med` using 2 as the argument. If you want, you can combine this step with the one above by adding the call to `.round()` to the end of the statement.
 4. Print `med`.
 5. Return `med`.
3. Create a dataframe called `hh` by using `pd.read_csv()` to read "households.csv". Use the keyword `index_col='id'` to set the index to the `id` field for each household.
4. Create a dataframe called `q` by using `pd.read_csv()` to read "quantities.csv". Again use the keyword `index_col='id'` to set the index to the `id` field for each household.
5. Determine the income quintile of each household by setting `hh` column "quint" to the result of calling the Pandas function `pd.qcut()` on the household income data. The `qcut()` function divides its input into bins and returns a series containing the bin number of each input record. The first argument to `qcut()` should be the column of incomes, `hh['inc']`, the second argument should be 5 to request quintiles, and the third should be `labels=[1,2,3,4,5]` to have the quintiles labeled 1-5. Please note that `pd.qcut()` is a standalone Pandas function like `pd.read_csv()`: it's not a series or dataframe method.
6. As in the earlier exercise, create a variable called `pd1` that is equal to 53.35 and one called `pd2` equal to 55.27. Then create a variable called `dp` that is equal to `pd2 - pd1`.
7. Compute the ETRs by multiplying 100 times `dp` times the `qd2` column of `q` divided by the `inc` column of `hh`. Store the result in the `hh` dataframe as column 'etr'. Notice that it's not necessary for the quantity and income data to be in the same dataframe. Also, this would work correctly no matter what order `hh` and `q` were in because Pandas will use the indexes to match up the income and quantity variables.
8. Now aggregate and print the ETRs by quintile alone by setting `med_q` to the result of calling `print_groups()` with arguments `hh` and `['quint']`.

9. Next, aggregate and print the results by type alone by setting `med_t` to the result of calling `print_groups()` with arguments `hh` and `['type']`.
10. Finally, aggregate and print the results by both type and quintile by setting `med_b` to the result of calling `print_groups()` with arguments `hh` and `['type', 'quint']`.
11. Print the index for `med_b`. Notice that it's a list of tuples with the type as the first element and the quintile as the second element.
12. Print an appropriate heading and then list the detailed medians for type 3 by printing `med_b[3]`. The `[3]` instructs Pandas to pick out all the elements where 3 is the first element in the index tuple. See the Tips section for an alternative approach.
13. Print an appropriate heading and then list the medians for the 5th quintile by printing `med_b[:,5]`. The 5 indicates that tuples with 5 in the quintile position should be printed, and the colon is a placeholder that indicates that all types (the first element in the tuples) should be included.
14. Print an appropriate heading and then print the median ETR for type 3, quintile 5, by printing `med_b[3,5]`.
15. Finally, to emphasize the power of the automatic alignment built into Pandas, print an appropriate heading and then print `med_b - med_b[:,1]`. That will show how much higher the ETR is for each quintile for each type relative to the first one for that type. Pandas will automatically align the data to ensure that the household type matches when doing the subtraction.

Submitting

Once you're happy with everything and have committed all of the changes to your local repository, please push the changes to GitHub. At that point, you're done: you have submitted your answer.

Tips

- The last few steps printing out subsets of the results rely heavily on the order of the values in the square brackets matching the order of the groups in the index. An alternative approach that is a little more verbose but avoids relying on the ordering is to use the `.xs()` method (short for cross-section). It allows the specific part of the index (here, the types or the quintiles) to be specified explicitly. We'll use `.xs()` in future exercises; for now, here are some quick examples showing how it would be used here:

```
# printing type 3:
print( med_b.xs(3,level='type') )

# printing quintile 5:
print( med_b.xs(5,level='quint') )
```